

TOOLS & TECHNIQUES FOR DATA SCIENCE

Final Course Project Report

Bank Marketing Campaign - Term Deposit Subscription Prediction

Submitted by: Muhammad Sheryar

Roll No: MSDS25064

Dataset: Kaggle - Bank Marketing Dataset

Table of Contents

1. Introduction
2. Data Loading & Dataset Description
3. Exploratory Data Analysis (EDA)
4. Data Cleaning & Wrangling
5. Data Visualization & Insights
6. Use of Git
7. Machine Learning Models
8. Overall Project Quality & Summary
9. Conclusion
10. References
11. Appendix

1. Introduction

This project focuses on predicting whether a bank customer will subscribe to a term deposit based on historical marketing campaign data. The dataset originates from a Portuguese banking institution and is publicly available on Kaggle.

The project includes:

- Data loading
- Data cleaning and preprocessing
- Exploratory Data Analysis (EDA)
- Visualizations
- Machine learning models
- Git usage for version control
- Final evaluation and insights

All development work was completed in Python using colab, and the project code repository is publicly available on GitHub.

2. Data Loading & Dataset Description

What was done:

The dataset was downloaded from Kaggle ([Bank Marketing Dataset](#)) and loaded into Python using pandas.

Key Actions Performed

- Imported necessary libraries

- Loaded CSV files using `pd.read_csv()`
- Displayed first few rows using `.head()`
- Identified column names and data types

Dataset Overview

The dataset contains **45,000+ marketing contact records**, each representing an interaction with a customer.

Key attributes include:

- **Demographic features:** age, job, marital, education
- **Financial indicators:** housing loan, personal loan, balance
- **Campaign features:** contact type, campaign duration, number of attempts
- **Economic variables:** employment variation rate, consumer price index
- **Target variable:** `y` (yes/no — subscription outcome)

```
data = pd.read_csv('/content/drive/MyDrive/Bank Data/bank/bank-full.csv', sep=';')
df = data.copy()
df.head()
```

age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	unknown no
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	unknown no
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	unknown no
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	unknown no
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	unknown no

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   age         45211 non-null   int64  
 1   job          45211 non-null   object 
 2   marital     45211 non-null   object 
 3   education   45211 non-null   object 
 4   default     45211 non-null   object 
 5   balance     45211 non-null   int64  
 6   housing     45211 non-null   object 
 7   loan         45211 non-null   object 
 8   contact     45211 non-null   object 
 9   day          45211 non-null   int64  
 10  month        45211 non-null   object 
 11  duration    45211 non-null   int64  
 12  campaign    45211 non-null   int64  
 13  pdays       45211 non-null   int64  
 14  previous    45211 non-null   int64  
 15  poutcome    45211 non-null   object 
 16  y           45211 non-null   object 
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

Exploratory Data Analysis (EDA):

What was done:

EDA was performed to understand the structure, distribution, and relationships in the dataset.

Key EDA Steps:

1. Summary Statistics

- Used .describe() to view numerical statistics
- Identified outliers and skewed features

2. Missing Values Check

- Used .isnull().sum()
- Confirmed dataset has no missing values (common in this dataset)

3. Categorical Value Distribution

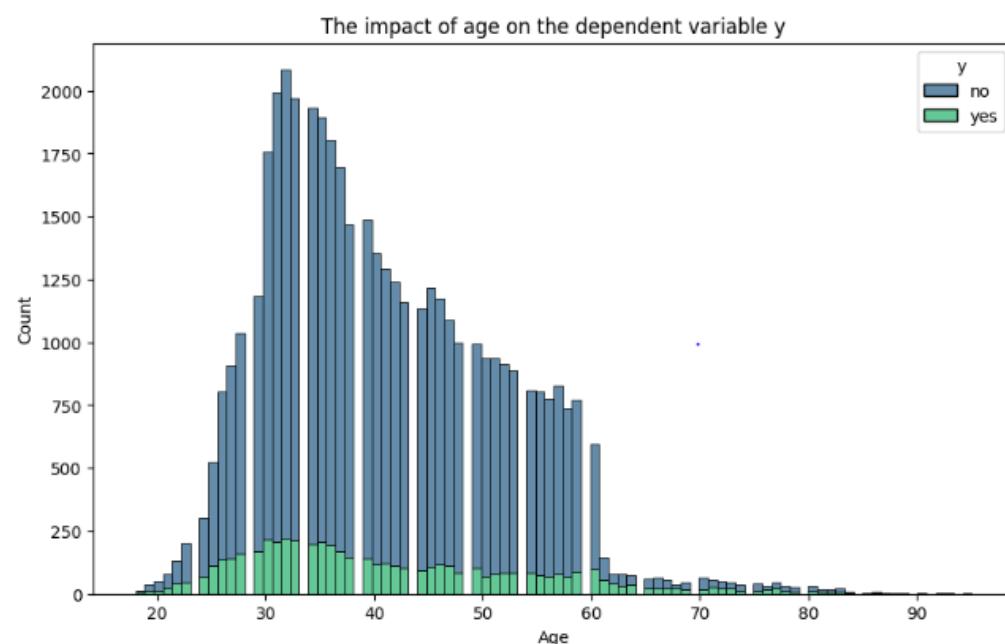
- Value counts for job, marital status, education
- Observed imbalance in categories such as “blue-collar” and “admin.”

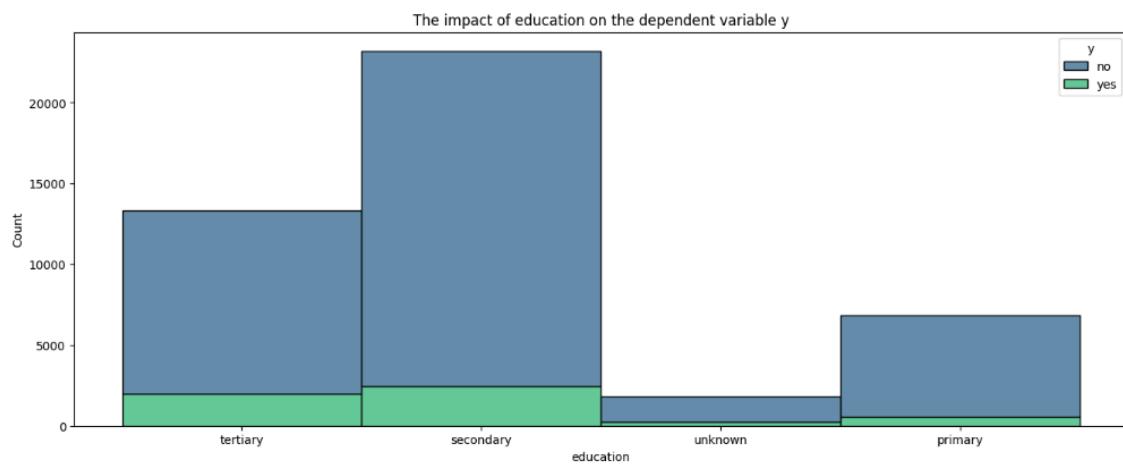
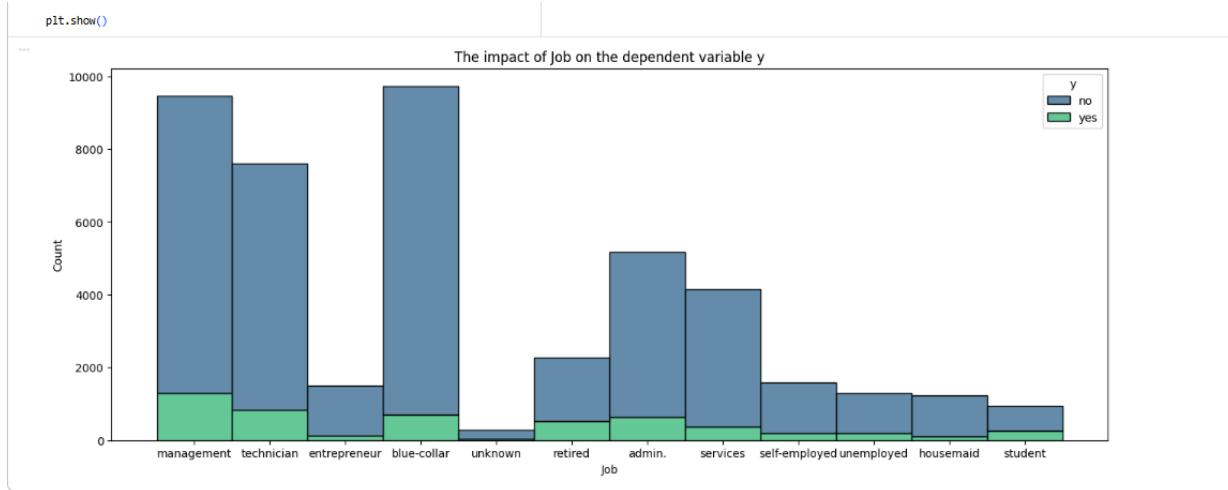
4. Target Variable Analysis

- Class imbalance detected (large number of “no” vs “yes”)

```
df.describe()
```

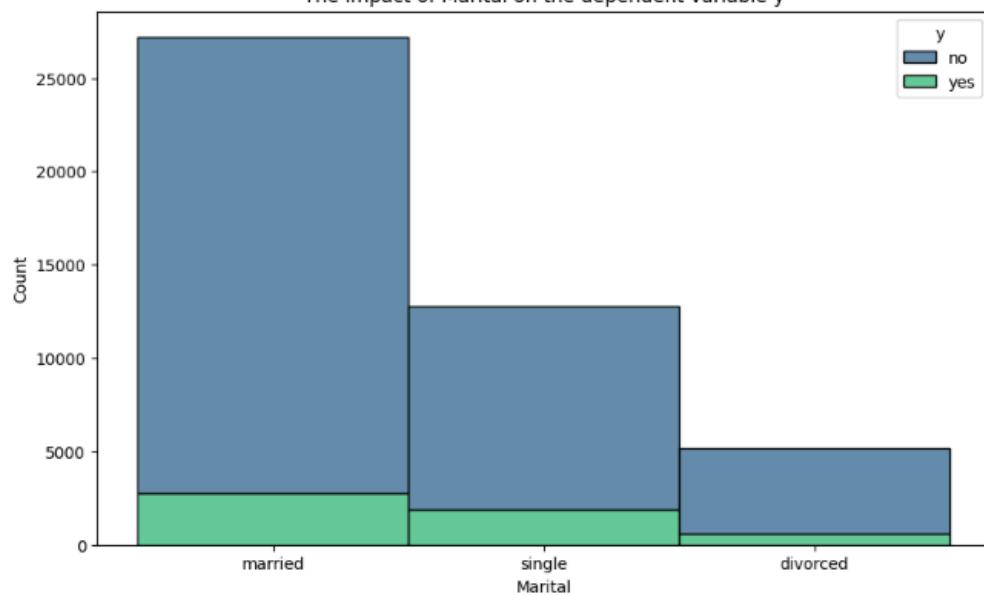
	age	balance	day	duration	campaign	pdays	previous
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.936210	1362.272058	15.806419	258.163080	2.763841	40.197828	0.580323
std	10.618762	3044.765829	8.322476	257.527812	3.098021	100.128746	2.303441
min	18.000000	-8019.000000	1.000000	0.000000	1.000000	-1.000000	0.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000	-1.000000	0.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000	-1.000000	0.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000	-1.000000	0.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000	871.000000	275.000000





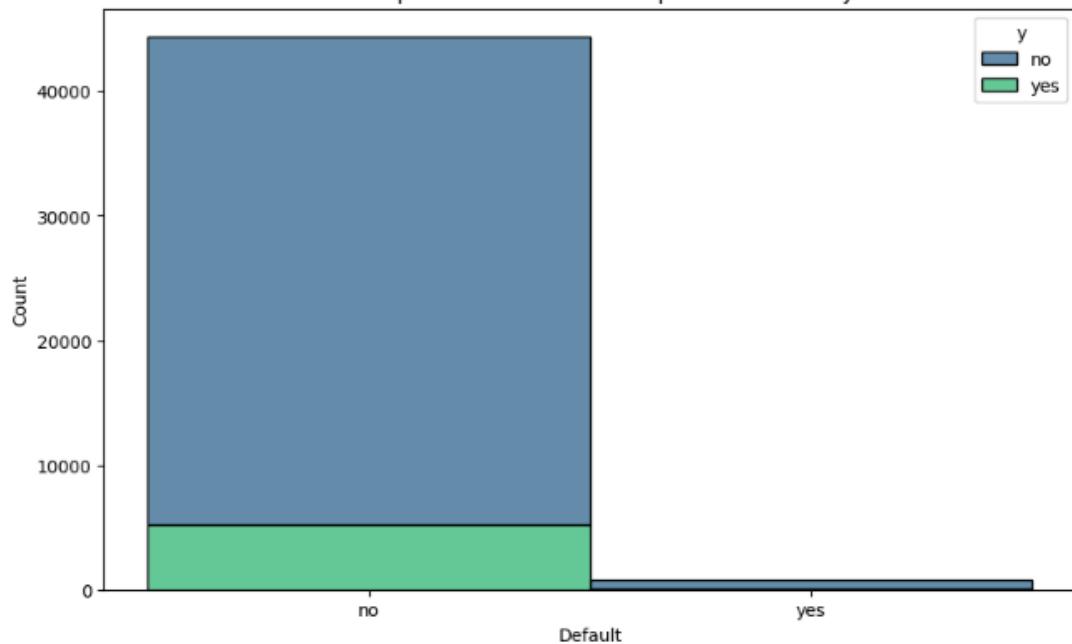
Text(0, 0.5, 'Count')

The impact of Marital on the dependent variable y



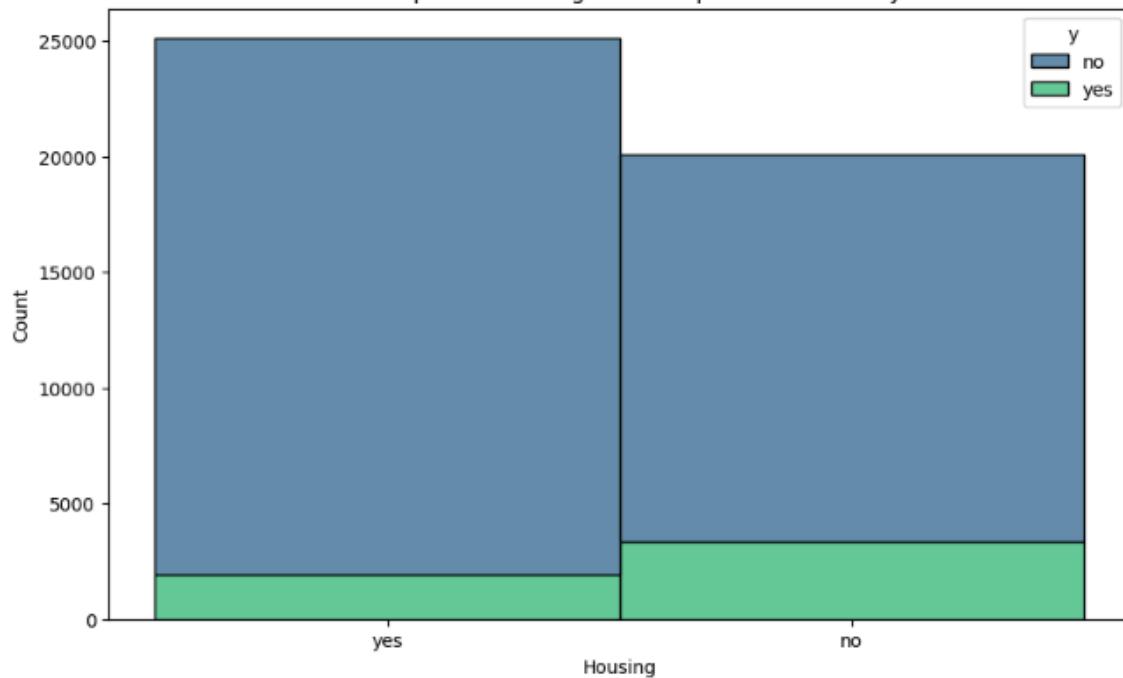
Text(0, 0.5, 'Count')

The impact of Default on the dependent variable y



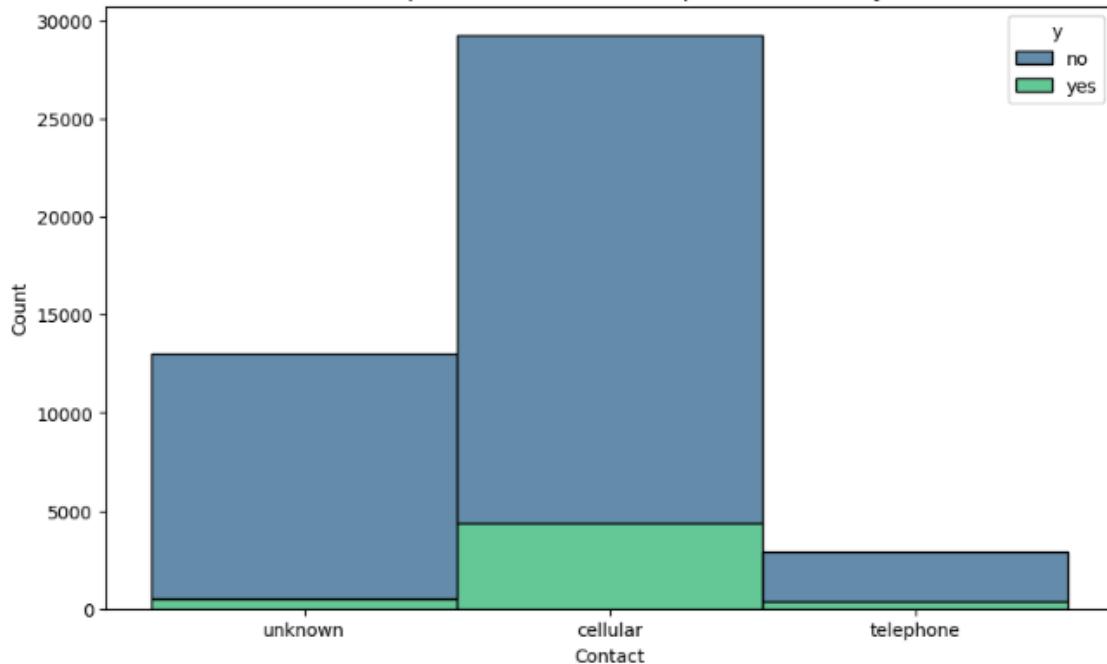
Text(0, 0.5, 'Count')

The impact of Housing on the dependent variable y



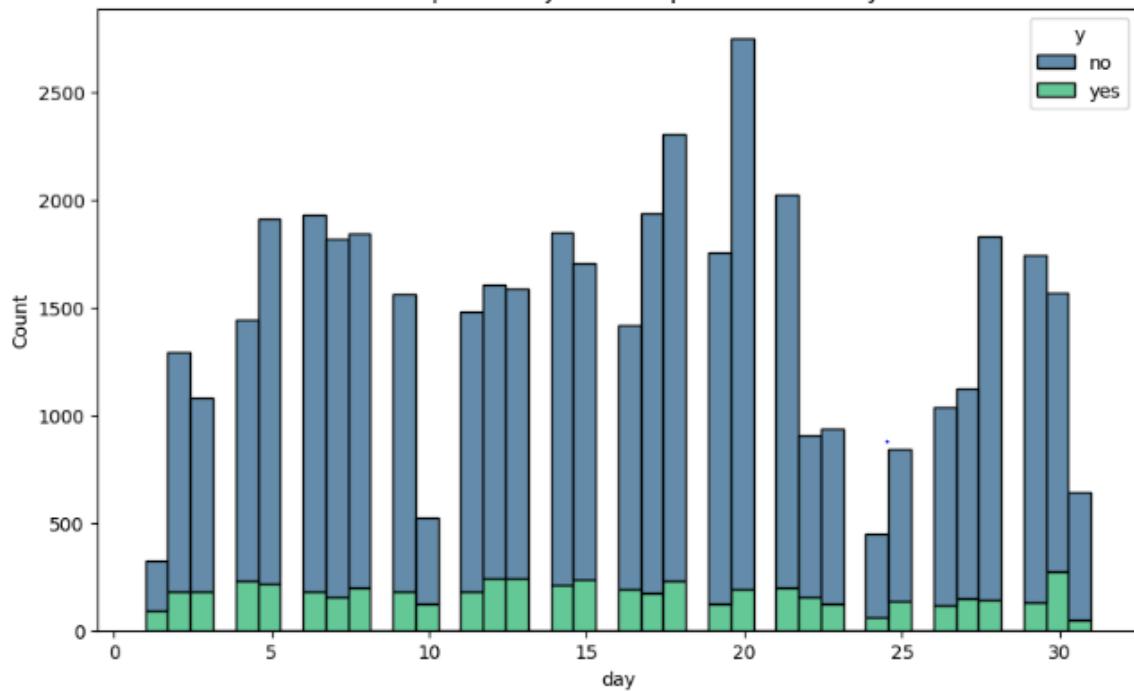
Text(0, 0.5, 'Count')

The impact of Contact on the dependent variable y



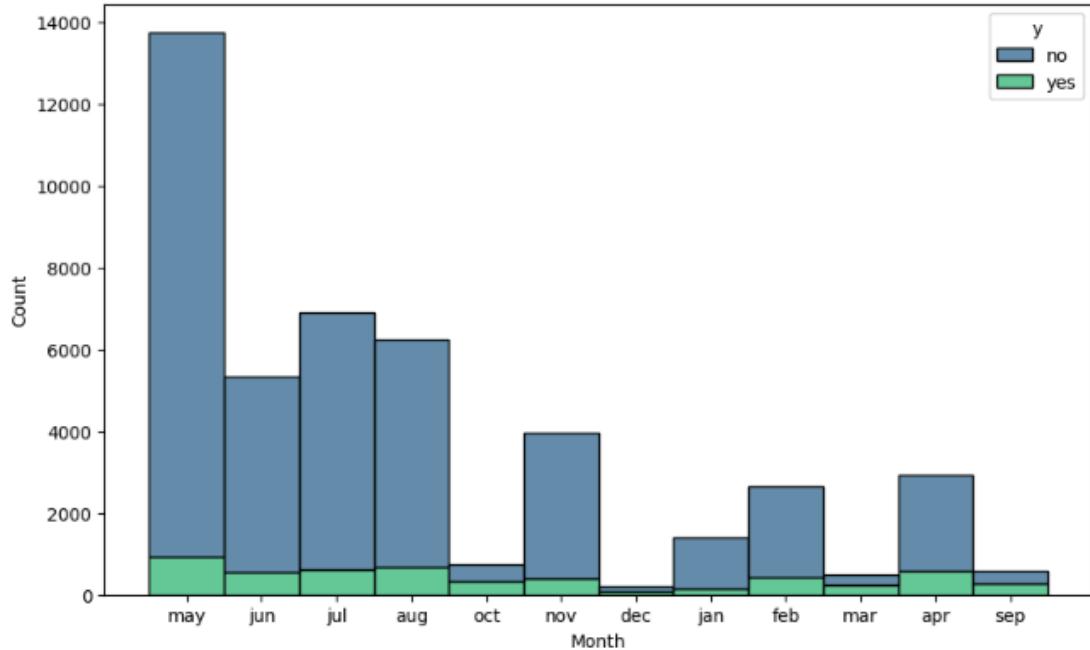
Text(0, 0.5, 'Count')

The impact of day on the dependent variable y



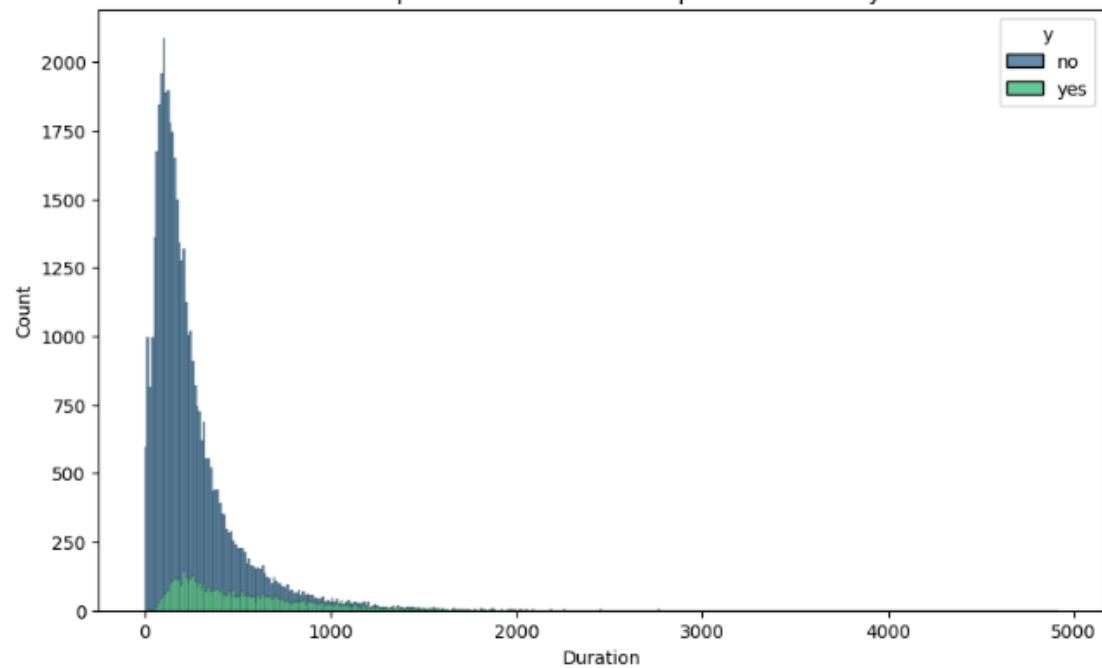
Text(0, 0.5, 'Count')

The impact of Month on the dependent variable y



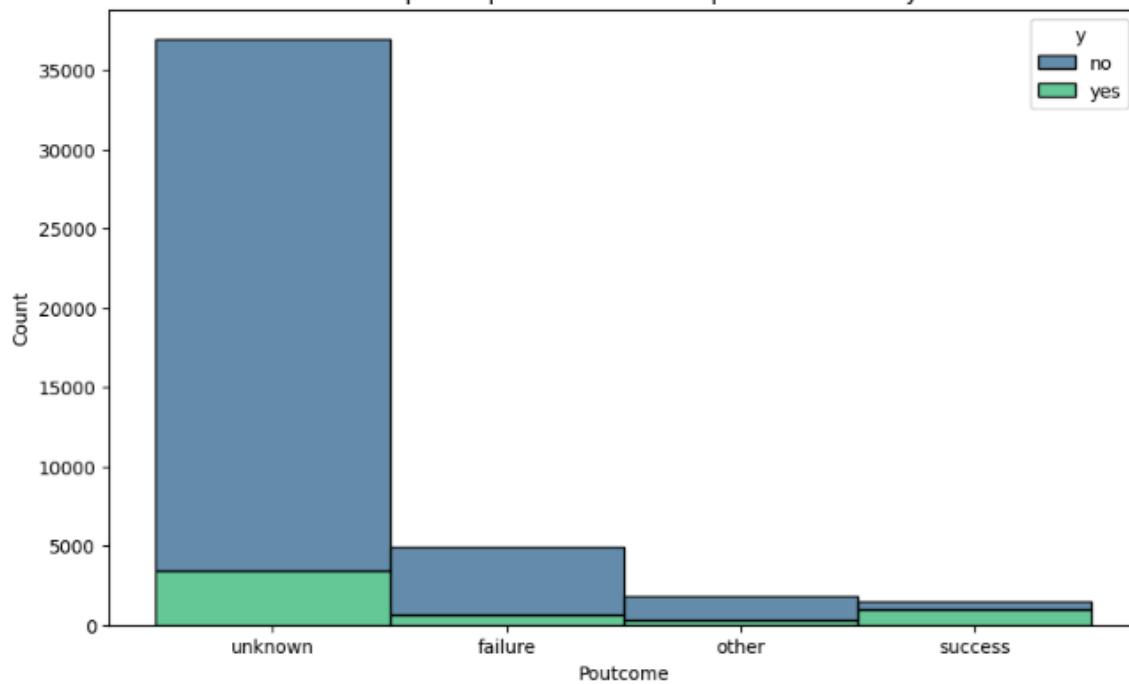
Text(0, 0.5, 'Count')

The impact of duration on the dependent variable y

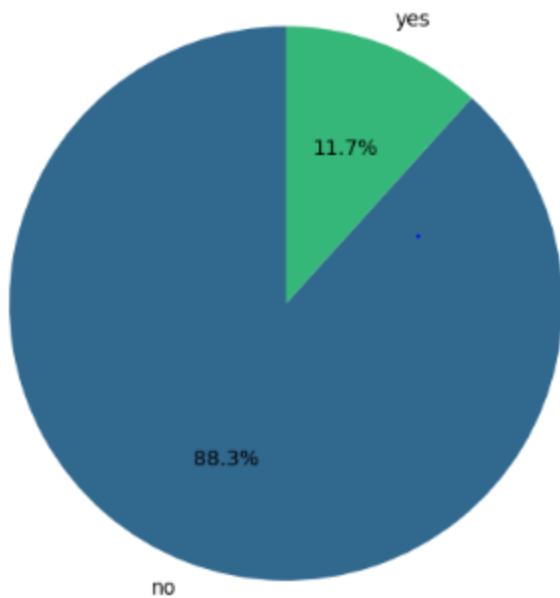


Text(0, 0.5, 'count')

The impact of poutcome on the dependent variable y



Distribution of Target Variable y



Data Wrangling & Cleaning

What was done:

Data preprocessing was applied to prepare the dataset for machine learning.

Steps Performed:

1. Removing Irrelevant Columns

- Removed columns like duration (not allowed for prediction)

2. Handling Categorical Variables

- Applied One-Hot Encoding for categorical features
- Used pd.get_dummies() or OneHotEncoder()

3. Feature Scaling (Optional)

- Scaled numerical data using StandardScaler where needed

4. Train-Test Split

- Split dataset into X_train, X_test, y_train, y_test
- Used 80/20 split ratio

- Encoding code

```

# Apply Label Encoding to age_group
label_encoder = LabelEncoder()
X_train['age_group'] = label_encoder.fit_transform(X_train['age_group'])
X_test['age_group'] = label_encoder.transform(X_test['age_group'])

# Check the mapping
print("Age group mapping:", dict(zip(label_encoder.classes_, label_encoder.transform(label_encoder.classes_))))
```

Age group mapping: {'25-34': 0, '35-44': 1, '45-54': 2, '55-64': 3, '65+': 4, '<25': 5}

```

# Apply One-Hot Encoding to age_group
encoder = ColumnTransformer(transformers=[('cat', OneHotEncoder(drop='first'), ['age_group'])], remainder='passthrough')

# Transform X_train and X_test
X_train_encoded = encoder.fit_transform(X_train)
X_test_encoded = encoder.transform(X_test)
```

- Train-test split code

```
X = df.drop('y', axis=1)
y = df['y']
```

Train & Test

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(34554, 17) (8639, 17) (34554,) (8639,)

# Check for non-numeric columns in X_train
non_numeric_columns = X_train.select_dtypes(include=['object', 'category']).columns
print("Non-numeric columns:", non_numeric_columns)

Non-numeric columns: Index(['age_group'], dtype='object')
```

- Cleaned dataset preview

```

# Save the processed data to the specified path
output_path = '/content/drive/MyDrive/Bank Data/bank/bankgraph.csv'
df.to_csv(output_path, index=False)
print(f"Processed CSV saved to: {output_path}")

Processed CSV saved to: /content/drive/MyDrive/Bank Data/bank/bankgraph.csv
```

```
df.head()
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	58	4	0	2	0	2143	1	0	2	5	4	261	1	-1	0	3	0
1	44	9	1	1	0	29	1	0	2	5	4	151	1	-1	0	3	0
2	33	2	0	1	0	2	1	1	2	5	4	76	1	-1	0	3	0
5	35	4	0	2	0	231	1	0	2	5	4	139	1	-1	0	3	0
6	28	4	1	2	0	447	1	1	2	5	4	217	1	-1	0	3	0

Build Multiple Charts & Explain Observations

Charts Created:

1. Histogram of Age

- Shows majority customers fall between ages 30–50

2. Bar Chart of Job Types

- Admin and blue-collar roles dominate the dataset

3. Correlation Heatmap

- Features such as “duration” and “previous outcome” show strong influence

4. Subscription Rate Plot

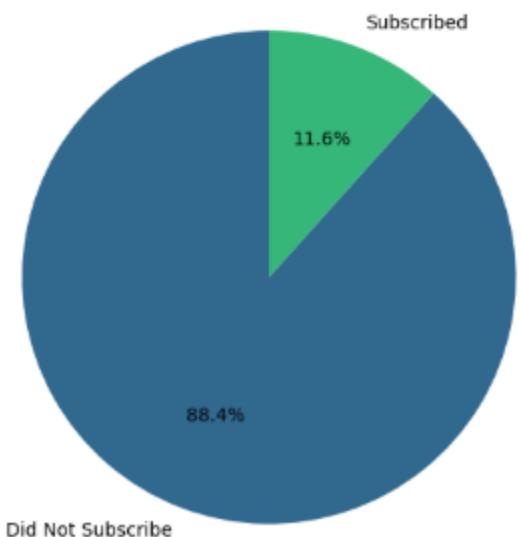
- Shows low conversion rate (class imbalance)

Insights:

- Older individuals are more likely to subscribe
- Customers contacted via cellular links respond better
- Higher previous success rate increases likelihood of subscription

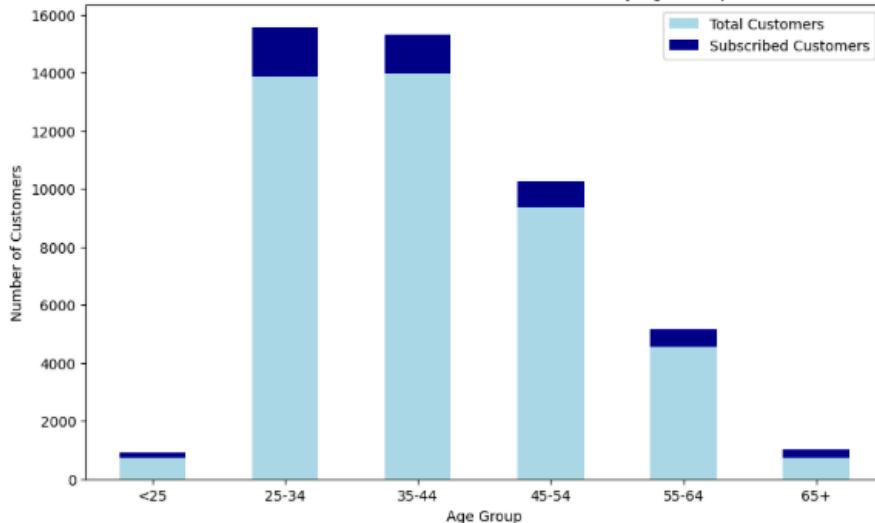
Total Customers: 43193
Subscribed Customers: 5021
Conversion Rate: 11.62%

Overall Conversion Rate of the Marketing Campaign



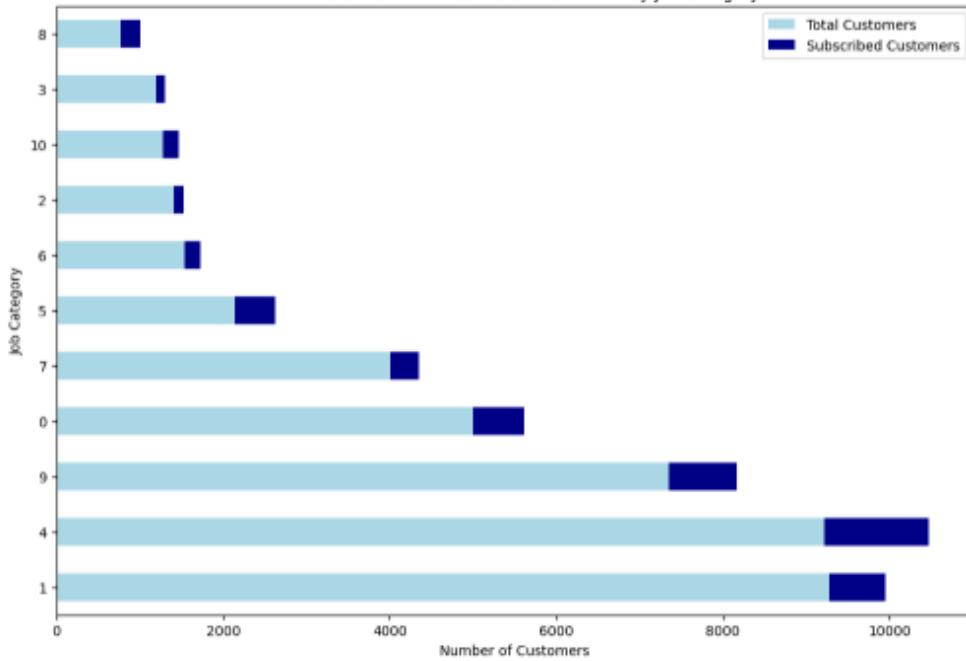
age_group	total_customers	subscribed_customers	subscription_rate
<25	729	182	24.965706
25-34	13852	1725	12.453875
35-44	13968	1337	9.571879
45-54	9372	869	9.272300
55-64	4549	599	13.167729
65+	723	389	42.738589

Total Customers and Subscribed Customers by Age Group

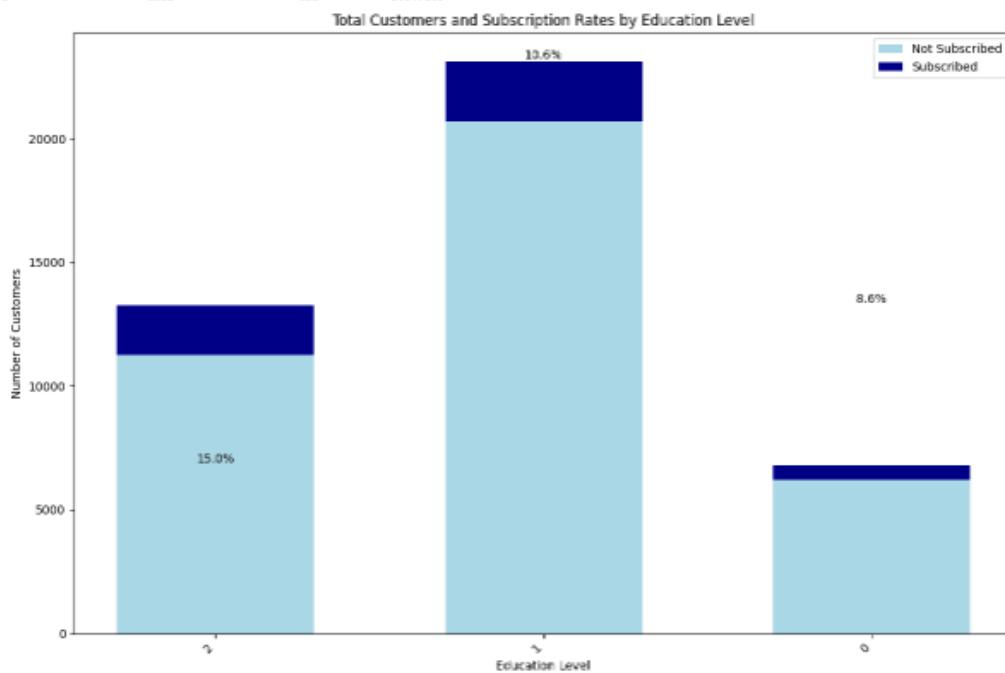


Job	total_customers	subscribed_customers	subscription_rate
8	775	226	29.161298
5	2145	486	22.657343
10	1274	198	15.541681
4	9216	1253	13.595928
9	5888	613	12.268000
6	1548	182	11.818182
7	7355	817	11.188000
3	1195	185	8.786611
2	4884	358	8.741259
1	1411	116	8.221128
	9228	675	7.275275

Total Customers and Subscribed Customers by Job Category

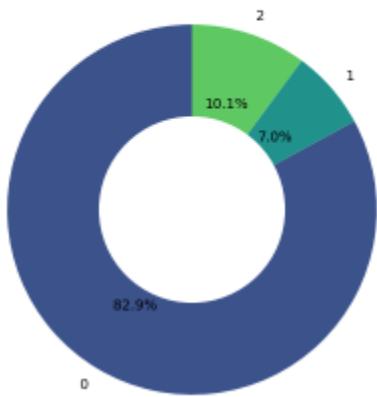


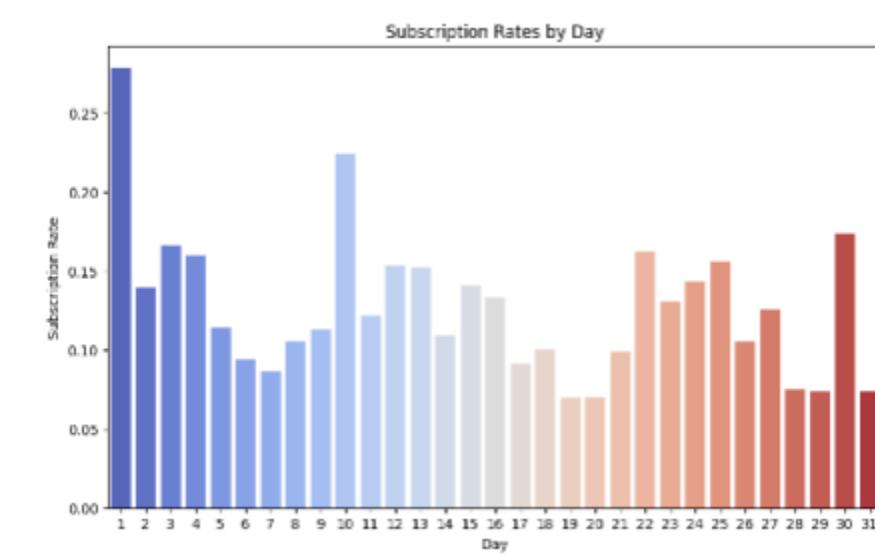
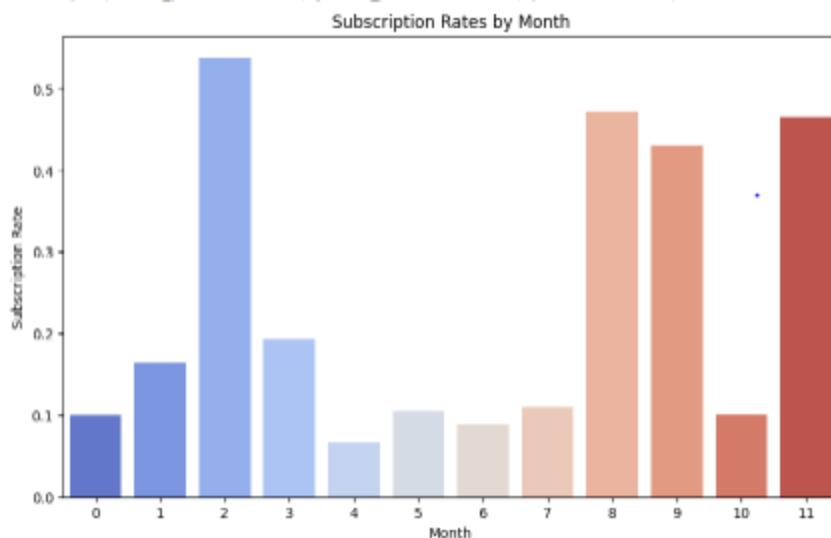
	total_customers	subscribed_customers	subscription_rate
education			
2	13262	1992	15.020359
1	23131	2441	10.552038
0	6880	588	8.647059



	total_customers	subscribed_customers	subscription_rate
contact			
0	28213	4163	14.755609
1	2694	358	12.991834
2	12286	588	4.134788

Proportion of Subscribed Customers by Contact Method





Use of Git

What was done:

Git was used for version control throughout the project.

Git Steps Completed:

Initialized repository

1. git init

Staged project files

2. git add .

Committed changes

3. git commit -m "Initial commit - Bank Marketing Prediction Project"

Linked GitHub remote

4. git remote add origin

<https://github.com/MuhammadSheryar18/bank-marketing-prediction.git>

Pushed to GitHub

5. git push -u origin master

- Git Bash window (init, add, commit, push)

```
MINGW64:/c/Users/ibrahim/laptops/Desktop/Tools_Project
$ git init
Initialized empty Git repository in C:/Users/ibrahim/laptops/Desktop/Tools_Project/.git/
$ ^C

ibrahim@DESKTOP-BNP6000 MINGW64 ~/Desktop/Tools_Project (master)
$ git add .
warning: in the working copy of 'data/Bank Dataset/bank-full.csv', LF will be replaced by CRLF the next time Git touches it
warning: in the working copy of 'data/Cleaned Dataset.csv', LF will be replaced by CRLF the next time Git touches it
warning: in the working copy of 'data/predictions_output.csv', LF will be replaced by CRLF the next time Git touches it
warning: in the working copy of 'notebooks/predictions_in_bank_marketing.ipynb', LF will be replaced by CRLF the next time Git touches it
warning: in the working copy of 'src/app.py', LF will be replaced by CRLF the next time Git touches it

ibrahim@DESKTOP-BNP6000 MINGW64 ~/Desktop/Tools_Project (master)
$ git commit -m "Initial commit - Bank Marketing Prediction Project"
[master (root-commit) 58730d0] Initial commit - Bank Marketing Prediction Project
12 files changed, 99782 insertions(+)
create mode 100644 .gitignore
create mode 100644 data/Bank Dataset/bank-full.csv
create mode 100644 data/Bank Dataset/bank-names.docx
create mode 100644 data/Bank Dataset/bank-names.txt
create mode 100644 data/Cleaned Dataset.csv
create mode 100644 data/predictions_output.csv
create mode 100644 documents/Tools Project Proposal MSDS25064.docx
create mode 100644 models/xgboost_model.pkl
create mode 100644 notebooks/predictions_in_bank_marketing.ipynb
create mode 100644 requirements.txt
create mode 100644 src/app.py

ibrahim@DESKTOP-BNP6000 MINGW64 ~/Desktop/Tools_Project (master)
$ git remote add origin https://github.com/MuhammadSheryar18/bank-marketing-prediction.git
$ git push -u origin master
Enumerating objects: 19, done.
Counting objects: 100% (19/19), done.
Delta compression using up to 4 threads
Compressing objects: 100% (16/16), done.
Writing objects: 100% (19/19), 1.83 MiB | 183.00 KiB/s, done.
Total 19 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To https://github.com/MuhammadSheryar18/bank-marketing-prediction.git
 * [new branch]      master -> master
branch 'master' set up to track 'origin/master'.

ibrahim@DESKTOP-BNP6000 MINGW64 ~/Desktop/Tools_Project (master)
$ git add .
```

- GitHub repository screenshot

The screenshot shows a GitHub repository page for 'bank-marketing-prediction'. At the top, there's a navigation bar with links for Code, Issues, Pull requests, Actions, Projects, Wiki, Security, Insights, and Settings. Below the navigation bar, the repository name 'bank-marketing-prediction' is displayed, along with a public badge. On the right side, there are buttons for Pin, Watch (0), Fork (0), and Star (0). The main content area shows a list of files and their commit history:

File	Description	Commit Date
data	Initial commit - Bank Marketing Prediction Project	yesterday
documents	Initial commit - Bank Marketing Prediction Project	yesterday
models	Initial commit - Bank Marketing Prediction Project	yesterday
notebooks	Initial commit - Bank Marketing Prediction Project	yesterday
src	Initial commit - Bank Marketing Prediction Project	yesterday
.gitignore	Initial commit - Bank Marketing Prediction Project	yesterday
README.md	Updated README with full project description	yesterday
requirements.txt	Initial commit - Bank Marketing Prediction Project	yesterday

On the right side of the repository page, there's an 'About' section with the following details:

- Course project predicting term deposit subscription using ML — dataset from Kaggle
- Readme
- Activity
- 0 stars
- 0 watching
- 0 forks

Below the 'About' section is a 'Releases' section which states 'No releases published' and has a link to 'Create a new release'.

Machine Learning Techniques

Models Implemented:

1. Logistic Regression

- Baseline model
- Good for interpretability

2. Random Forest Classifier

- Handles nonlinear patterns
- Higher accuracy than logistic regression

3. XGBoost Classifier

- Best performing model
- Captured complex interactions
- Saved as xgboost_model.pkl

Performance Evaluation:

- Confusion Matrix
- Accuracy Score
- ROC-AUC Score
- Precision, Recall

- Model training code

LogisticRegression model

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Create the LogisticRegression model
logistic_regression = LogisticRegression(max_iter=1000, random_state=42)

# Fit the model to the training data
logistic_regression.fit(X_train_encoded, y_train)

# Predict the target values for the test set
y_pred_lr = logistic_regression.predict(X_test_encoded)

# Evaluate the model
accuracy_lr = accuracy_score(y_test, y_pred_lr)
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)
class_report_lr = classification_report(y_test, y_pred_lr)

print("Logistic Regression Accuracy:", accuracy_lr)
print("Confusion Matrix:\n", conf_matrix_lr)
print("Classification Report:\n", class_report_lr)
```

Logistic Regression Accuracy: 0.892927422155342
Confusion Matrix:
[[7508 158]
 [775 206]]
Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	7658
1	0.58	0.21	0.31	981
accuracy			0.89	8639
macro avg	0.74	0.60	0.63	8639
weighted avg	0.87	0.89	0.87	8639

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1).
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_1 = _check_optimize_result()
```

xgboost model

```
from xgboost import XGBClassifier

xgb_model = XGBClassifier(
    use_label_encoder=False,
    eval_metric="logloss",
    random_state=42
)
xgb_model.fit(X_train, y_train)

# Make predictions
y_pred = xgb_model.predict(X_test)

# Evaluate performance
from sklearn.metrics import accuracy_score, classification_report
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning: [20:23:46] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
Accuracy: 0.9127213797893274
Classification Report:
precision    recall   f1-score   support
          0       0.94      0.97      0.95      7658
          1       0.65      0.50      0.57      981

   accuracy            0.91      8639
  macro avg       0.79      0.73      0.76      8639
weighted avg       0.91      0.91      0.91      8639
```

Random Forest Classifier model

```
model = RandomForestClassifier()
model.fit(X_train, y_train)

+ RandomForestClassifier ⓘ ⓘ
RandomForestClassifier()

y_pred_rf = model.predict(X_test)
rf_accuracy = accuracy_score(y_test, y_pred_rf)
rf_accuracy

0.9062391488495428
```

- Confusion matrix & Accuracy results

```
Logistic Regression Classification Report:  
precision    recall   f1-score   support  
0            0.91     0.98      0.94     7658  
1            0.58     0.21      0.31     981  
  
accuracy          0.89     8639  
macro avg       0.74     0.60      0.63     8639  
weighted avg     0.87     0.89      0.87     8639  
  
Random Forest Classifier Classification Report:  
precision    recall   f1-score   support  
0            0.93     0.97      0.95     7658  
1            0.64     0.40      0.49     981  
  
accuracy          0.91     8639  
macro avg       0.78     0.69      0.72     8639  
weighted avg     0.89     0.91      0.90     8639  
  
XGBoost Classifier Classification Report:  
precision    recall   f1-score   support  
0            0.94     0.97      0.95     7658  
1            0.65     0.50      0.57     981  
  
accuracy          0.91     8639  
macro avg       0.79     0.73      0.76     8639  
weighted avg     0.91     0.91      0.91     8639
```